Neel Gali

Generative AI Engineer

neelgali7@gmail.com | +1(475)-291-1455 | Austin, TX

Professional Experience

AI Engineer *Jul 2025 – Present*

Epik Solutions | Enterprise Security AI

- Architected bespoke **multi-agent systems** with **LangGraph** and **Google ADK** using ReAct patterns and hierarchical delegation, moving beyond prototypes to ship directly inside live enterprise environments.
- Drove prototype-to-production transitions for 5+ agentic customer workflows, resolving state-management, data-readiness, and security perimeter blockers preventing enterprise-grade AI, cutting go-live time by 40%.
- Engineered **TypeScript** and Python **MCP servers** deployed on **Google Cloud Run**, forming connective tissue between **Vertex AI** and customer legacy APIs, data silos, and security perimeters.
- Built **evaluation pipelines** and observability frameworks tracking LLM-native metrics—tokens/sec, cost-per-request, p95 latency—ensuring every agentic deployment met accuracy, safety, and latency SLAs.
- Surfaced 12+ recurring friction patterns across customer AI deployments, converting real-world field insights into formal **Google Cloud** product feature requests that directly shaped Engineering roadmaps.
- Embedded with 3 enterprise engineering teams, instilling Google-grade eval-driven development, **CI/CD**, and observability practices that improved long-term AI adoption and reduced support escalations by 35%.
- Implemented **granular tracing** with **LangSmith** and state-management optimization across agent graphs, reducing debugging cycles by 45% for customer-facing AI pipelines across all production accounts.

AI/ML Engineer *Jul 2024 – Jun 2025*

Samay Systems LLC | B2B SaaS Compliance

- Applied **OpenAI GPT-4** prompt engineering, structured output, and self-reflection techniques for compliance summarization, cutting post-processing errors by 45% over rule-based extraction on 10,000+ monthly documents.
- Built stateful **multi-agent** compliance workflows with **CrewAI** and **LangGraph** in 2024, enabling parallel document processing and reducing end-to-end pipeline latency by 35% at production scale.
- Architected **RAG pipelines** pairing **Azure AI Search** vector retrieval with LLM generation, achieving 89% retrieval precision and powering context-aware compliance reports for 10,000+ enterprise users.
- Fine-tuned **HuggingFace BERT** and T5 for compliance document classification, processing 10,000+ monthly documents and reducing manual auditor effort by 70% across enterprise clients.
- Provisioned AI services on **AWS** (SageMaker, ECS, Lambda) with **Terraform** IaC, enforcing CI/CD model versioning and cutting deployment errors by 40% across all enterprise compliance tenants.
- Exposed AI outputs via **FastAPI** streaming endpoints consumed by **React.js** dashboards, improving enterprise user task-completion rates by 30% and reducing support escalations across 10,000+ production users.

SDE Intern *Jun 2023 – Aug 2023*

Monarch Tractor | Autonomous Vehicles

- Built **Python** ML preprocessing and inference pipelines transforming raw sensor and image data into training-ready datasets, directly supporting 3 production computer vision systems for autonomous vehicle navigation.
- Deployed **AWS Lambda** pipelines with **Docker** containers and **GitHub Actions** CI/CD, reducing ML pipeline failures by 40% and accelerating model iteration cycles across all production computer vision systems.

Associate ML Engineer *Jul 2021 – Aug 2022*

Zemoso Technologies | LabOps Platform

- Trained **PyTorch** classification and regression models on scientific experiment datasets, building **scikit-learn** feature engineering pipelines that reduced model training cycle time by 30% across LabOps modules.
- Fine-tuned **HuggingFace BERT**-based NLP models for domain-specific document classification and named entity extraction over lab reports, achieving 85%+ accuracy and replacing manual data entry for 500+ scientists.
- Built **AWS Lambda** and **S3**-triggered data ingestion pipelines for automated ML training dataset refresh, reducing data prep effort by 40% and ensuring model freshness across 3 production computer vision systems.
- Surfaced ML model outputs via **Spring Boot** REST APIs and **React.js** front-ends, enabling accurate experiment tracking for 500+ scientists across 10+ regulated scientific workflows.

Skills

Programming Languages: Python, TypeScript, JavaScript, Java, SQL, Bash
Agentic AI & LLMs: LangGraph, CrewAI, Google ADK, MCP Servers, OpenAI API (GPT-4o), Vertex AI
ML & NLP: PyTorch, HuggingFace Transformers, BERT, T5, Model Fine-tuning, Prompt Engineering, RAG
Vector Databases & Search: Vertex AI Search, Pinecone, OpenSearch, Semantic Search, Embedding Models
Cloud & Infrastructure: GCP (Vertex AI, Cloud Run, GCS), AWS (SageMaker, Lambda, ECS, S3), Terraform
Observability & Evaluation: LangSmith, Evaluation Pipelines, Granular Tracing, LLM-native Metrics (tokens/sec, cost-per-request)
Frameworks & Tools: FastAPI, Spring Boot, React.js, Git, JIRA, Agile/Scrum, RESTful API Design, Docker, Kubernetes

Education

Master's in Computer Science, Oklahoma State University, Stillwater
Bachelor's in Computer Science, JNTU Hyderabad, India