Neel Gali

AI Engineer

neelgali7@gmail.com | +1(475)-291-1455 | Austin, TX

Professional Experience

AI Engineer *Jun 2025 - Present*

Epik Solutions

- Built LLM inference and RAG pipelines using Python and LangChain on enterprise security event data, substantially reducing analyst task completion time by 40% in production.
- Implemented model drift detection and performance monitoring using AWS CloudWatch and custom Python metrics, ensuring production LLM accuracy stayed above SLA thresholds across all clients.
- Containerized and deployed AI inference services on AWS Kubernetes using Docker and Terraform, supporting automated CI/CD pipelines that maintained 99.9% uptime across production.
- Provided proactive L1/L2 AI platform support via Slack channels and weekly office hours, resolving integration blockers and onboarding enterprise users across 15+ client organizations efficiently.
- Partnered with product, security, and infrastructure stakeholders to define AI system requirements, delivering Python and Java production API features reliably across multiple Agile sprint cycles.

AI Engineer *Jul 2024 - Jun 2025*

Samay Systems LLC

- Engineered production LLM inference and text summarization pipelines on AWS Bedrock for enterprise compliance, processing 10K+ documents per month and cutting auditor effort by 70%.
- Implemented RAG-style querying over compliance corpora using Pinecone and OpenSearch vector databases, enabling natural language queries for multiple enterprise clients, boosting accuracy to 89%.
- Architected multi-tenant compliance SaaS platform on AWS, deploying highly containerized Python microservices with Docker and Kubernetes, maintaining strict tenant isolation, scaling to 10,000+ users.
- Implemented LLM evaluation frameworks to detect and remediate model drift in compliance logic, achieving 95% accuracy on evidence summarization and control scoring across enterprise deployments.
- Built Python and Java data ingestion pipelines with AWS Step Functions and Lambda supporting CI/CD automation across compliance control collections for 15+ enterprise client environments.

SDE Intern *Jun 2023 - Aug 2023*

Monarch Tractor

- Built Python-based ML preprocessing and inference pipelines transforming raw sensor and image data into training-ready datasets supporting 3 production computer vision systems for autonomous vehicles.
- Built AWS Lambda data pipelines with Docker containers and GitHub Actions CI/CD, enabling consistent ML training and reducing pipeline failures by 40% across production systems.

Software Engineer *Jul 2020 - Aug 2022*

ZeMoSo Technologies

- Delivered full-stack features for a regulated scientific cloud platform using React and async Spring Boot Java REST APIs, ensuring accurate data tracking for 500+ scientists.
- Designed RESTful API contracts and asynchronous service patterns across 10+ core product modules, significantly improving overall team engineering standards and production system reliability by 35%.
- Authored 100+ JUnit test cases and resolved critical production issues, integrating CI/CD pipelines to enforce quality checks and maintain data integrity in a regulated environment.

Skills

Programming Languages: Python, Java, JavaScript, TypeScript, SQL, Bash
AI & LLMs: OpenAI API, Anthropic API, HuggingFace, LangChain, RAG, Prompt Engineering, Inference Pipelines, Model Drift Detection
Vector Databases & Search: Pinecone, OpenSearch, Azure AI Search
Cloud & DevOps: AWS (Bedrock, Lambda, S3, EC2, Step Functions), Docker, Kubernetes, Terraform, GitHub Actions, CI/CD
Frameworks & Databases: Spring Boot, FastAPI, React.js, PostgreSQL, MongoDB, Redis

Education

Master's in Computer Science, Oklahoma State University, Stillwater
Bachelor's in Computer Science, JNTU Hyderabad, India