Lohit Gali

Senior Data Scientist

lohitgali@gmail.com | +1(475)-291-1455 | San Francisco, CA

PROFESSIONAL SUMMARY

Senior Data Scientist with 5+ years of experience transforming complex business challenges into actionable ML problem statements. Proficient in Python and SQL for robust feature engineering and model development across classification, regression, and time series forecasting. Proven expertise in Deep Learning (PyTorch, TensorFlow/ Keras) and MLOps for building production-aware, scalable solutions, driving measurable business impact.

TECHNICAL SKILLS

Programming & Data Query: Python, SQL, Pandas, NumPy, PySpark

Machine Learning & Deep Learning: Classification, Regression, Time Series Forecasting, Clustering, Anomaly Detection, Gradient Boosting (XGBoost, LightGBM, CatBoost), Neural Networks, PyTorch, TensorFlow/Keras, Scikit-learn, Feature Importance (SHAP)

MLOps & Cloud Platforms: Databricks (Spark, Delta Lake, MLflow), AWS, Azure, GCP, Docker, Kubernetes, Git, GitHub Actions, Airflow, Jira, Model Monitoring, Reproducible Training

Data Engineering & Analytics: ETL, Data Pipelines, Data Warehousing, Data Lakes, Delta Lake, Apache Kafka, Power BI, Tableau, Statistics (Hypothesis Testing, Experiment Design, A/B Testing)

Agentic AI Development: LangChain, LLMs (OpenAI, Anthropic), Prompt Engineering, Tool-Calling Agents, RAG

PROFESSIONAL EXPERIENCE

Senior Data Scientist | Epik Solutions

Jul 2025 - Present

• Led the end-to-end development of a real-time anomaly detection system using Deep Learning (PyTorch) and time series forecasting models, reducing critical security incident detection time by 70% and minimizing potential financial losses by $5M annually for compliance teams.

• Architected and implemented scalable MLOps pipelines on Azure Databricks with MLflow for model versioning, automated evaluation, and monitoring for drift/performance, improving model retraining efficiency by 45% for financial risk models.

• Translated complex financial regulations into clear ML problem statements, designing and deploying gradient boosting (XGBoost) and neural nets for structured data for classification and regression tasks, achieving 92% prediction accuracy for compliance risk.

• Developed and integrated Agentic development workflows using Python and LangChain to automate data profiling and error analysis across large financial datasets, enhancing data quality assurance and reducing manual effort by 60% for data engineering teams.

• Designed and conducted rigorous experiment design and A/B testing for new feature rollouts, leveraging statistics (hypothesis testing, confidence intervals) to validate uplift and inform strategic product decisions, resulting in a 15% increase in user engagement.

• Optimized feature engineering pipelines on Databricks Spark for tabular data, processing petabytes of data for classification and regression models, reducing feature generation time by 30% and improving model performance by 8% AUC for research analysts.

• Collaborated with engineering teams to package models for batch scoring pipelines and real-time APIs using Docker and Kubernetes on AWS, ensuring reproducibility and seamless integration into production systems for global stakeholders.

Data Scientist | Samay Systems LLC

Jul 2024 - Jun 2025

• Developed and deployed supervised learning models (LightGBM, CatBoost) for classification and regression to predict compliance risks, improving prediction accuracy by 18% and reducing false positives by 25% for security analysts.

• Implemented robust feature engineering pipelines using Python and SQL to extract and transform

compliance-related data from diverse sources (AWS, Azure, GitHub), supporting downstream ML model development.

• Designed and validated time series forecasting models to predict audit readiness and resource allocation, achieving a MAPE of 10% and optimizing operational planning by 20% for project managers.

• Leveraged deep learning techniques (TensorFlow/Keras) for document classification and entity recognition from unstructured text data, automating evidence collection by 35% and accelerating audit preparation.

• Contributed to MLOps initiatives by implementing automated evaluation and testing frameworks for production models, ensuring model reliability and performance consistency in a SaaS environment.

Machine Learning Engineer | Arvam Tech

Jul 2023 - May 2024

• Engineered data ingestion and transformation pipelines using Python and Databricks to consolidate asset and operational datasets into analytics-ready formats, enabling accurate trend analysis and forecasting for enterprise asset management.

• Developed and deployed clustering algorithms (k-means, DBSCAN) to segment asset usage patterns, identifying anomalies and optimizing maintenance schedules, leading to a 15% reduction in unplanned downtime.

• Built event-driven workflows to process streaming updates from IoT devices and ensure near-real-time availability of analytical data, supporting proactive risk identification and operational decision-making.

• Designed and implemented model monitoring dashboards using Power BI, visualizing key performance indicators and model predictions for stakeholders, ensuring transparency and enabling rapid response to performance degradation.

Associate Data Scientist | RHYM Tech

Aug 2022 - May 2023

• Developed analytical models using Python and SQL to analyze high-volume cybersecurity telemetry and event data, identifying emerging threats and optimizing detection rules for the threat intelligence platform.

• Performed extensive data cleaning, transformation, and feature extraction to support downstream ML tasks, improving the quality and completeness of training data by 30% for anomaly detection models.

• Applied time-series aggregation and anomaly detection logic to identify emerging trends across security telemetry, enhancing threat visibility and reducing false positive alerts by 20%.

• Created interactive analytical dashboards using Python libraries to visualize distributions, trends, and risk indicators, effectively communicating insights to both technical and non-technical stakeholders.

Associate Developer | ZeMoSo Technologies

Jul 2021 - Aug 2022

• Developed robust data ingestion scripts using Python and SQL to process high-volume experimental data from biotech instruments, ensuring data quality and availability for scientific research teams.

• Implemented foundational data cleaning and transformation pipelines to standardize raw sensor and operational data, enabling initial exploratory data analysis for future machine learning model development.

• Assisted in the design and optimization of relational database schemas for operational efficiency, improving data retrieval speeds by 20% for analytical reporting requirements in a lab automation context.

• Created interactive data visualizations using Python libraries to communicate key performance indicators and trends for lab operations, supporting rapid decision-making and stakeholder reporting.

EDUCATION

Master’s in Computer Science | Oklahoma State University, Stillwater, OK

Bachelor’s in Computer Science | JNTU Hyderabad, India